

Alex Chen

Email: alex.chen@example.com | Phone: 06 12345678 | Location: Amsterdam, Netherlands

LinkedIn: linkedin.com/in/alex-chen-example

20 January 2026

Dear Hiring Team,

Example Corp's Q4 platform report named job latency as a recurring customer pain point. At Nexus Labs I led the migration from a monolithic job runner to queue-based workers on Kubernetes, cutting p95 latency from twelve seconds to three. That is the same reliability work your Amsterdam platform squad describes in the posting.

I also built the internal deployment CLI your team would extend: forty engineers use it daily, and staging deploy time dropped from twenty-five minutes to six. The CLI wraps Helm releases, environment promotion, and smoke checks so platform changes stay repeatable. Earlier at DataFlow I owned ingestion on-call for two years and wrote runbooks the support team still uses when pipelines stall overnight.

On the API side I designed public REST v2 with explicit versioning and rate limits, which let three enterprise customers integrate without breaking existing contracts. I paired that work with PostgreSQL query tuning that cut warehouse sync time by sixty per cent. I am comfortable translating product goals into SLOs, Grafana dashboards, and alert routes — the observability culture your listing emphasises.

Before Nexus I spent four years at DataFlow building Python ETL services at two million events per day. That period taught me how to document incidents clearly and how to keep migrations reversible when customers depend on daily freshness.

I am based in Amsterdam and can work on site as required. Example Corp's focus on mid-market logistics customers matches the B2B platform work I have shipped for the past four years. Thank you for considering my application.

Kind regards,

Alex Chen